

The present work was submitted to:

Chair of Process and Data Science
RWTH Aachen University

State-Based Process Monitoring in Traditional Event Logs for Concept Drift Detection

Bachelor Thesis

presented by

Bora Deren
438377

First Examiner: Prof. Dr. Wil M. P. van der Aalst
Second Examiner: Prof. Dr. Sander Leemans

Advisor: Dr. Alessandro Berti

Aachen, October 5, 2026

Acknowledgments

If you want to incorporate it, keep in mind that, according to the official regulations, *independent work* is a very important aspect of writing a thesis. Thus, you can be grateful for the emotional support by your favorite pet, yet we recommend you to be careful if you want to write anything more specific. In general, we suggest to omit this section.

Abstract

The abstract should be a concise summary of what you have done. It should cover, in roughly 1 to 2 sentences per topic:

- The domain, i.e., explain the application domain in which your thesis is applicable.
- The problem, i.e., motivate what the problem is you are solving.
- The method, i.e., describe the method you have proposed, and, compare it, briefly & high level to other approaches. In particular, focus on the benefits of your approach versus the other approach(es)
- The evaluation, i.e., how did you evaluate your method, what results did you obtain and what do the results indicate?

Contents

1	Introduction	1
1.1	Motivation	2
1.2	Problem Statement	5
1.3	Research Questions	6
1.4	Research Goals	7
1.5	Contributions	8
1.6	Thesis Structure	8
2	Related Work	11
2.1	Drift Detection in Control-Flow	11
2.2	Drift Detection Beyond Control-Flow	12
2.3	Clustering-Based Drift Detection	12
2.4	States in Process Mining	12
2.5	Event Logs with Known Drifts	13
2.6	Positioning of This Thesis	14
3	Preliminaries	15
4	Rheon	17
5	Intra-Case State	19
5.1	Feature Extraction	19
5.2	Dimensionality Reduction	19
5.3	State Clustering	19
5.4	Drift Signals	19
6	Resource State	21
6.1	Feature Extraction	21
6.2	Dimensionality Reduction	21
6.3	State Clustering	21
6.4	Drift Signals	21
7	Inter-Case State	23
7.1	Feature Extraction	23
7.2	Dimensionality Reduction	23
7.3	State Clustering	23
7.4	Drift Signals	23

Contents

8	Implementation	25
8.1	Kairo	25
8.2	Dashboard	25
9	Evaluation	27
9.1	Experimental Setup	27
9.2	Results	27
9.3	Threats to Validity	28
10	Discussion	29
11	Conclusion	31
A	Some additional stuff	33

List of Figures

1.1	A waiting-time drift that control-flow does not show	4
1.2	Drifts on two different time scales	6
2.1	Transition system with states for remaining time prediction	13
4.1	Placeholder figure.	17

List of Tables

4.1 Placeholder table.	17
--------------------------------	----

1 Introduction

First of all, we encourage you to read Prof. van der Aalst’s guide on how to write papers [19].

In this section, you introduce the work to your audience. Even though this is the first chapter of your thesis, you should not necessarily write this chapter first. It is often easier to write this chapter at the end of the thesis writing. In this part of your introduction, i.e., the part before the Motivation section (Section 1.1), you write a paragraph (not just a sentence) for each of the bullet points mentioned in the abstract:

- Domain, i.e., explain the application domain in which your thesis is applicable.
- Problem, i.e., motivate what the problem is you are solving.
- Method, i.e., describe the method you have proposed, and, compare it, briefly & high level to other approaches. In particular, focus on the benefits of your approach versus the other approach(es)
- Evaluation, i.e., how did you evaluate your method and what do the results indicate?

Optionally, you can finalize this part with a little overview of what is still to come in this chapter: “The remainder of this chapter is structured as follows. In Section 1.1, we motivate the need for ... In Section 1.2, we provide a formal problem statement. ...”

Before we dive into the motivation section, a few general guidelines for writing:

1. Write *active* and in the *present tense*.
2. Avoid “optionality” as much as possible, i.e., usage of the following words should (pun intended) be minimized:
 - could
 - would
 - should
 - may
 - might
 - can
 - will
 - ...

1 Introduction

Good: We present a method that ...

Bad: We will present a method that ...

Good: We use function f and compute its inverse.

Bad: We can use the function f and will compute its inverse.

3. Do not use `\\` to generate line breaks, but an empty line
4. Position the caption of a Figure below the figure.
5. Position the caption of a Table above the table.
6. Write Section Headings and Titles Like This
Good: Approximation Bias in Unstable Systems
Bad: Approximation bias in unstable systems
7. Be consistent in your references.
 - Make sure that the name of the same author is always the same (using DBLP helps for this)
 - Titles should be consistent. Either use the style previously described, i.e., Approximation Bias in Unstable Systems or Approximation bias in unstable systems (this is allowed in the references, opposed to your own section headings and titles, however, *BE CONSISTENT*).

1.1 Motivation

Played-out event logs do not change, but real processes do. Following van der Aalst [18], the situation in which a process is changing while being analyzed is called a concept drift. For example, two activities are concurrent at the start of an event log but then become sequential. According to van der Aalst [18], typical triggers of concept drift are seasonal effects, such as higher demand in December, and changed market conditions. Classical process mining techniques mostly assume that the process at the start and at the end of the log is the same [6]. If a drift occurs, process models or key performance indicators (KPIs) computed over the entire log mix two different processes. The process owner must therefore detect whether and when the process changed, and what changed.

A drift does not restrict itself to the order of activities. Processes run in a context, but process analysis sometimes ignores it. As described by van der Aalst [18], process context covers how cases influence each other like competing for the same resources. A single case is therefore not sufficient to explain the observed behavior alone. To explain the observed behavior and its changes, we analyze dependencies within cases, between cases, and with resources. For these dependencies, Senderovich et al. [17] introduce the terms intra-case and inter-case dependencies. This thesis adopts this terminology and considers a process from three perspectives:

- *Intra-case perspective*: The history of a single case up to a specific time, covering which activities happened in which order and how often. Drift example: A new trace variant dominates.

- *Resource perspective:* The working style of resources, such as their workload, service times, handovers, and the activities they perform. Drift example: A resource no longer exists and other resources take over the work.
- *Inter-case perspective:* The current situation of all running cases of the process, covering arrivals, waiting times, and case attributes. Drift example: The mean waiting time between consecutive events of a case increases.

From all three perspectives concept drifts can take place, but they can also influence each other. For example, a new trace variant with a frequent activity becomes dominant (intra-case perspective). The resource responsible for that activity then has more work to do (resource perspective), which affects the waiting times in the entire process (inter-case perspective).

We describe each drift by three properties:

- *Drift type:* What changes, e.g., the control-flow, the reassignment of an activity to another resource, or the waiting time between consecutive events.
- *Drift perspective:* The perspective to which the drift type belongs. For example, control-flow drifts belong to the intra-case perspective, reassignment drifts to the resource perspective, and waiting time drifts to the inter-case perspective.
- *Drift mode:* How the change takes effect over time, e.g., as a sudden, gradual, or recurring change.

In the literature, the terms drift type, drift mode, and drift perspective are used interchangeably. We reserve the term drift type for what changes, the term drift mode for how it changes, and the term drift perspective for the group a drift type belongs to.

Existing concept drift detection approaches cover this picture only partially:

- *Control-flow focus:* Most of the work on drift detection focuses on control-flow. Out of the 38 drift detection papers in the survey of Sato et al. [16], only six focus on the time perspective, three on the data perspective, and none on the resource perspective. In our framework, waiting times and case attributes fall under the inter-case perspective, and service times under the resource perspective. A concept drift that only changes the time or resource perspective stays undetected with control-flow approaches.
- *Single KPIs:* Approaches beyond control-flow often track single KPIs. Adams et al. [3] turn each KPI chosen by the user into a time series over time windows and detect drifts in each series separately. A drift therefore shows up as separate change points in several KPIs. This approach does not deliver a compact description of the overall process situation. For steady-state detection, Kraus et al. [11] show that combining complementary process characteristics is more accurate than relying on a single one.

1 Introduction




Figure 1.1: Visualization of the synthetic Rheon log generated with a mean waiting time drift. At the drift point (dashed), the mean waiting time between consecutive events of a case rises from 14.2 to 45 min, while all cases still follow process tree T . The inter-case state of each 12-hour window switches exactly there from A to an alternation of B and C. State D marks the run-out of the log after the last case arrivals. The alternation of B and C is drawn schematically.

- *Drift modes:* How a change takes effect over time is also important information for the process owner. Most of the existing approaches focus on sudden drifts only [16].
- *Missing test data:* Sato et al. [16] argue that real-world logs are unsuitable for measuring detection accuracy, since their drift points are unknown. They find only three public synthetic datasets with injected drifts. These datasets cover only the control-flow and time perspectives and do not mix different drift modes. None of them contains a resource drift or a change in case arrivals. A recent benchmark by Padella and Vinci [15] injects drifts into five real logs, but again only into the control-flow and activity durations.

Figure 1.1 shows this gap using a synthetic event log generated by Rheon. Rheon is our generator for synthetic event logs with injected concept drifts in all three perspectives (Chapter 4). This log has 300,299 cases with 1,946,540 events, seven activities, and seven resources. Its cases arrive between January 1 and December 31, 2020. On April 19, 2020, the mean waiting time between consecutive events of a case jumps from 14.2 minutes to 45 minutes. The activities, the control-flow, the resources, and the dominant resource of each activity remain unchanged. A detector that only compares relations between activities sees the same control-flow before and after the drift point and misses the change.

The state-based approach of this thesis aims to detect drifts in all three perspectives and to show what changed. Instead of separate change points per KPI, our states combine many features into one compact description. Rheon provides test logs with known drift points in all three perspectives. In Berti et al. [5], execution states from object-centric event logs are computed for objects. We adapt the approach from objects

to cases in our case-centric framework and define our intra-case states. As an extension to their work, our approach also takes into account resource and inter-case perspectives and adapts the idea to them.

Overall, we compute states from an event log by extracting features, compressing, and clustering. The input of our approach is an event log, and the output consists of intra-case, resource, and inter-case states whose changes over time indicate drifts. In Figure 1.1 one can visually see the general idea of our approach for the inter-case perspective. We describe each 12-hour window by features such as the number of active cases, arrivals, completions, and the mean waiting time between consecutive events of a case. Following the execution states of Berti et al., a principal component analysis (PCA) compresses these features. A 2×2 self-organizing map (SOM) is trained on these compressed feature vectors. Each of its four cells is one state, and each window gets the state of its best-matching cell.

The pipeline is unsupervised and does not see the drift point or the drift type. As seen in Figure 1.1, all 219 windows before the drift point fall into state A. From the drift point until December 31, the windows alternate between states B and C. The last seven windows, in which the log runs out, form state D. The first principal component explains 51% of the variance and is dominated by the mean waiting time between consecutive events of a case. It therefore also shows what changed: after the drift point, more time passes between the events of a case.

1.2 Problem Statement

An event log records a process over a long period, and the process often changes within this period. Such a concept drift affects the intra-case, resource, or inter-case perspective, or several of them at once. A detector that only captures control-flow drifts therefore misses drifts that only affect the resource or inter-case perspective.

Van der Aalst [18] names three main challenges in dealing with concept drifts:

1. *Change point detection*: Did the process change? If yes, when did it change?
2. *Change localization and characterization*: What has changed?
3. *Change process discovery*: How to capture and predict second-order dynamics?

According to van der Aalst, a process is not static even when there is no drift (first-order dynamics). This makes it even harder to detect changes of the process itself (second-order dynamics). Most approaches compare sliding windows of events with statistical methods. To localize drift points exactly and detect them at different time scales, the window length must vary. Moreover, drifts at different time scales can overlap, as in Figure 1.2.

This thesis focuses on challenges 1 and 2. To characterize a drift, our framework determines its drift perspective, its drift mode, and what exactly changed.

In conclusion, this thesis focuses on the following problem:

1 Introduction




Figure 1.2: A periodically changing process with drifts at two different time scales, taken from van der Aalst [18]. Shorter sliding windows are used to detect short-term drifts. The left half of each sliding window is compared with the right half to spot statistically relevant changes. The larger sliding windows at the bottom are used for long-term drifts.

- *Given:* An event log in which each event belongs to exactly one case. Each event carries a case ID, an activity, and a timestamp, and optionally a resource and case attributes. Drift points, drift types, drift perspectives, and drift modes are unknown.
- *Objective:* Compute intra-case, resource, and inter-case states whose changes over time indicate drifts. From these changes, detect the drift points and determine for each drift its perspective, its mode, and what changed.
- *Evaluation:* We compare the detected drift points with known drift points in two sets of logs. First, synthetic Rheon logs contain drifts in all three perspectives. Second, the benchmark of Padella and Vinci [15] injects sudden, gradual, and recurring control-flow and duration drifts into five real logs.

1.3 Research Questions

The main research question of this thesis is:

Can a process be monitored as a combination of intra-case, resource and inter-case states, and is this combination a useful indicator for a concept drift in traditional case-centric event logs?

We divide our main research question into five further questions:

- RQ1** *States:* How can we compute execution states from a traditional event log in the intra-case, resource and inter-case perspectives, and how interpretable are these states in explaining concept drifts?
- RQ2** *Detection:* Do the changes over time in those states signal drift points even if an event log has drifts in multiple perspectives?

- RQ3** *Characterization*: Do the states indicate the perspective and the mode of a drift, and which features suit which drift types?
- RQ4** *Clustering*: Which of the clustering algorithms SOM, k-means and Density-Based Spatial Clustering of Applications with Noise (DBSCAN) delivers the most robust and reliable states?
- RQ5** *Applicability*: In which processes and event logs is the state-based drift detection approach more reliable and when does it deliver deceptive results?

RQ1 addresses the first part of the main research question, monitoring a process through states. RQ2 to RQ5 cover the second part of the main research question, the usage of the states as an indicator for concept drift detection.

1.4 Research Goals

In order to answer the research questions, we follow seven research goals. Each goal lists the research questions it addresses in parentheses.

- RG1** *Terminology*: To clearly define drift type, drift mode and drift perspective (RQ3).
- RG2** *Test data*: Develop a generator for synthetic event logs that injects concept drifts of various drift types in two drift modes (sudden and gradual). Along with each log, the generator writes a metadata file that lists the injected drifts (RQ2, RQ3, RQ5).
- RG3** *Framework*: Develop a state-based framework for concept drift detection (RQ1, RQ2). It adapts execution states from Berti et al. [5] for objects to cases and extends them to resource and inter-case states.
- RG4** *Library*: Develop an AI-extended Python library that implements our framework (RQ1, RQ4). The library offers each result as a plot or text for large language models (LLM) and also provides LLM connections.
- RG5** *Dashboard*: Develop a user-friendly dashboard around the library (RQ1, RQ3). In the dashboard, users compute and analyze results like the states and try different pipelines and configurations with the help of an LLM-based assistant, the Copilot.
- RG6** *Evaluation*: Evaluate the robustness and flexibility of the framework against different processes, event logs and timespans (RQ2, RQ5)
- RG7** *Limitations*: Reveal the limitations of the approach and potential misleading results especially when using the Copilot (RQ5)

1.5 Contributions

This thesis has six main contributions. Each contribution lists its research goals and chapters in parentheses.

- *Terminology*: A terminology defining and clearly separating drift types (what changes), drift modes (how it changes over time), and drift perspectives (which perspective the drift type belongs to) (RG1).
- *Rheon*: (RG2, Chapter 4). A library that generates synthetic event logs and injects nine drift types from all three perspectives, each suddenly or gradually. For each generated log, Rheon creates a metadata file comprehensively describing how the log was generated and a list of drifts together with their drift points, types and metadata.
- *State-based framework*: A framework that computes the intra-case, resource and inter-case states (RG3, Chapters 5 to 7). The framework consists of four steps: feature extraction, dimensionality reduction with PCA, clustering into states, and drift signals. The drift signals compare state distributions of calendar windows with divergences, or state vectors with similarity measures.
- *kairo*: A Python library that implements the framework (RG4, Section 8.1). With the help of its user-friendly documentation, users create a pipeline configuration in kairo and run it. The library supports various features addressing different drift types, various scaling techniques, PCA, and SOM, k-means and DBSCAN for state clustering. It also computes state frequency distributions for different window sizes, offers multiple divergences and similarity measures, and provides various plotting options for the different pipeline steps.
- *Dashboard*: A dashboard built around kairo (RG5, Section 8.2). In a user-friendly user interface (UI), users can upload an event log, run the pipeline steps for each perspective and investigate the states in detail. With the LLM-based Copilot, users ask questions, gain insights about the results and request further steps.
- *Evaluation*: An evaluation on synthetic Rheon logs with different drift combinations from all perspectives and on the five real event logs of Padella and Vinci [15] (RG6, RG7, Chapters 9 and 10). The evaluation encompasses comprehensive comparisons between SOM, k-means and DBSCAN, and between different configuration parameters like scaling techniques, window sizes and feature selections.

1.6 Thesis Structure

The remainder of the thesis is structured as follows. Chapter 2 reviews the literature on concept drift detection and state-based process mining approaches and places this thesis within the literature. Chapter 3 introduces the notation for event logs, foundations of

concept drifts, PCA and clustering. Chapter 4 presents Rheon, which generates the synthetic event logs for the evaluation.

Chapters 5 to 7 define intra-case, resource and inter-case states and answer RQ1. Each of these chapters follows the same steps: feature extraction, dimensionality reduction, clustering, and drift signals. Chapter 8 describes the implementation of kairo and the dashboard tool.

Chapter 9 answers RQ2 to RQ5 using the Rheon logs and the benchmark logs of Padella and Vinci [15]. Chapter 10 discusses the design decisions and limitations of the framework. Chapter 11 summarizes the results and discusses future directions and potential extensions of the topic.

2 Related Work

This chapter gives an overview of the literature on different concept drift detection approaches, state-based process mining and positions the thesis in the literature. Section 2.1 presents detectors that focus on detecting control-flow drifts, and Section 2.2 focuses on the works on detectors beyond control-flow. Section 2.3 covers clustering-based detection approaches. Section 2.4 describes the role of states in process mining and reviews the state-based literature. In Section 2.5 we focus on publicly available event logs with known drifts as test data for the evaluation. Finally, Section 2.6 positions this thesis in the literature by explaining similarities and differences of our framework.

2.1 Drift Detection in Control-Flow

The most common strategy is a statistical test that compares the behavior of neighboring windows. In the survey of Sato et al. [16] 11 out of 36 papers use a statistical test approach. Bose et al. [6] apply statistical tests on follows relationships features between activities. The tool ProDrift developed by Maaradji et al. [14] uses a χ^2 test on the distribution of partially ordered runs, changes the window size adaptively and detects gradual drifts alongside sudden drifts. ProDrift receives higher F1-scores and detects drifts earlier than the approach of Bose et al. [6].

Other detectors represent the control-flow differently. In Tsinghua Process Concept Drift Detection (TPCDD) Zheng et al. [20] derive candidate drift points from a matrix of traces and relations and cluster them together using DBSCAN. With an error tolerance (ET) of 10 traces around the actual drift point, TPCDD achieves an F1-score of 1.0, whereas ProDrift needs an ET of about 100 traces for a comparable score on the same test data. CV4CDD-4D [12] learns to detect concept drifts using event logs with known sudden, gradual, incremental, and recurring drifts in an image of window similarities. CV4CDD-4D achieves an F1-score of 0.86 and finds the drift points more granular than the other approaches evaluated. It is also more robust against noise.

Beyond drift detection, Kraus and van der Aa [13] distinguish drift localization (what changed) from drift characterization (how it changed over time) and address the latter. In our terms, these correspond to drift types and drift modes.

All the detectors consider only the control-flow and therefore just the intra-case perspective of this thesis. Drifts that only change the waiting times, such as the one in Figure 1.1, stay hidden.

2.2 Drift Detection Beyond Control-Flow

Several detectors cover further perspectives. Brockhoff et al. [7] compare the trace distribution of two windows using the Earth Mover’s Distance (EMD) with regard to processing and sojourn times.

Three detectors detect drift points using the Pruned Exact Linear Time (PELT) algorithm on time series of features. Adams et al. [3] use control-flow, performance, data, resource, object features and test pairs of detected drifts for Granger causality. Klijn et al. [10] use features on resources and control-flow using an event knowledge graph (EKG). CVDrift by Khraiweh et al. [9] chooses the window size of each time series from its coefficient of variation (CV), with series for control-flow, time, and resources. On event logs with activity duration drifts CVDrift achieves F1-scores from 0.77 to 0.97.

PELT-based detectors observe a drift in many features but capture the inter-case perspective only through individual features such as arrival time or processing time.

2.3 Clustering-Based Drift Detection

Methodologically, the clustering-based detectors are the most similar to this thesis. Clustering-based detectors observe how the composition of clusters changes over time and newly emerging or vanishing clusters are a drift signal (Sato et al. [16]). Concept-Drift in Event Stream Framework (CDESf) developed by Barbon Junior et al. [4] clusters cases of an event stream using DBSCAN with respect to distances in control-flow and time. CDESf signals a drift when outlier cases become dense enough to form a new cluster. Version Clustering (Abarca Zúñiga et al. [1]) clusters windows according to their directly-follows frequencies into process versions using Hierarchical Density-Based Spatial Clustering of Applications with Noise (HDBSCAN). It signals drift points at the transitions between these versions. Using a tolerance of 5% of the cases, Version Clustering achieves F1-scores of 0.84 and 0.96 on two benchmarks.

Both approaches apply clustering in a feature space describing mostly the control-flow. Version Clustering (Abarca Zúñiga et al. [1]) applies clustering on full windows similar to resource and inter-case perspectives of this thesis, but in our thesis the states do not represent a process version.

2.4 States in Process Mining

Classically a state describes a running case for making predictions. van der Aalst [18] predicts remaining time of cases using annotated transition systems with states (Figure 2.1). Senderovich et al. [17] encode running cases for prediction in a state space from intra-case and inter-case dependencies.

Berti et al. [5] learn execution states of objects using sliding windows, PCA and SOM. In this thesis, we adapt the approach to cases and extend it to resource and inter-case perspectives beyond the intra-case perspective using calendar windows and feature extraction for concept drift detection.




Figure 2.1: Transition system annotated with time measurements, taken from van der Aalst [18]. Each node is a state of a running case and each arc is the start or completion of an activity. Each state collects the elapsed, remaining, and sojourn times of all past cases that visited it, as shown for state [p3,p4]. The gray path is a running case that reaches [p3,p4]. Its predicted remaining flow time is 42.56, the average remaining time observed in this state.

Kraus et al. [11] detect steady states from KPIs for each time window like the number of active cases. A fluctuating arrival rate is not seen as a concept drift by the paper whereas for this thesis it is a concept drift in the inter-case perspective.

2.5 Event Logs with Known Drifts

Adams et al. [2] compare eight detectors only on synthetic event logs with 132 control-flow drifts. None of the detectors performs well on all six evaluation goals, such as accuracy, latency, and robustness to noise. This kind of synthetic logs can be generated by CDLG [8]: It plays out process trees and injects drifts in the control-flow in four modes. The benchmark by Padella and Vinci [15], by contrast, has five real event logs, but with drifts only in control-flow and activity durations. None of these event logs has drifts in the resource perspective or arrival rates. Rheon follows the model of CDLG but adds multiple drift types from all three perspectives supporting sudden and gradual drifts (Chapter 4).

2.6 Positioning of This Thesis

Three gaps motivate the approach of this thesis:

- *Perspectives*: Most detectors detect only control-flow drifts. Detectors for further perspectives recognize drift points in individual features or entire trace distributions but do not describe the process situation compactly.
- *States*: Clustering-based detectors mainly cluster control-flow behavior. The state-based works in Section 2.4 use states for prediction, object analysis, or steady-state detection, but not for concept drift detection.
- *Test data*: None of the aforementioned works and tools comprise resource drifts or arrival rate drifts.

This thesis addresses all three gaps. Our states describe a process in all three perspectives separately (RQ1), and we evaluate if their changes signal drifts (RQ2). Rheon delivers test data with injected drifts in all three perspectives (RG2).

3 Preliminaries

In the preliminaries you present well-understood mathematical concepts that you need in your thesis. For example, you can define the natural numbers as $\mathbb{N} := \{1, 2, \dots\}$, and, correspondingly $\mathbb{N}_0 := \mathbb{N} \cup \{0\}$. A preliminary notion is either a well-defined commonly understood mathematical notion, e.g., sets, multisets, graphs, sequences, Petri nets, $\dots$, or, it is a concept clearly defined in another paper, i.e., you just adopt the notation (or a slight variation thereof). *Any concept you use should be defined in your thesis.* You should never write: “We use Workflow nets, a definition of these can be found here [X]”. If you use it, explain it.

Concepts that are unique to your approach are not part of the preliminaries, i.e., they are described in the approach section itself.

Some useful tips:

- When introducing a complex concept, use the following structure (always works):
 - Explain the concept informally.
 - Provide a formal definition of the concept.
 - Provide an example, using the formal definition, of the concept.

In your examples, try to be *as visual as you can*, often, an image says more than 5 pages of text.

- use commands, some of which are already provided in the preamble, e.g.,
`\newcommand*{\N}{\mathbb{N}}`

4 Rheon

Placeholder.

Table 4.1: Placeholder table.

Column A	Column B
Placeholder	Placeholder




Figure 4.1: Placeholder figure.

5 Intra-Case State

Placeholder.

5.1 Feature Extraction

Placeholder.

5.2 Dimensionality Reduction

Placeholder.

5.3 State Clustering

Placeholder.

5.4 Drift Signals

Placeholder.

6 Resource State

Placeholder.

6.1 Feature Extraction

Placeholder.

6.2 Dimensionality Reduction

Placeholder.

6.3 State Clustering

Placeholder.

6.4 Drift Signals

Placeholder.

7 Inter-Case State

Placeholder.

7.1 Feature Extraction

Placeholder.

7.2 Dimensionality Reduction

Placeholder.

7.3 State Clustering

Placeholder.

7.4 Drift Signals

Placeholder.

8 Implementation

Placeholder.

8.1 Kairo

Placeholder.

8.2 Dashboard

Placeholder.

9 Evaluation

In this chapter, you present your evaluation. You should have extensively discussed with your supervisor what you are evaluating and why ...

Some typical sections that you need here:

9.1 Experimental Setup

Describe the setup of your experiments. If you apply a pipeline of different techniques, show the pipeline. Present the parameters of your experiments in a table, and, describe them. Typical Elements:

- Input Data Set(s) Used
- Algorithm(s) Used
- Parameter(s) Used
- ...

9.2 Results

Show the results. Try to present your results as structured as you can. A good result (sub)section, follows the following structure:

- Describe what results you are going to Show
- Present an initial hypothesis about the results, e.g., We expect the quality metric M to behave like this, conditional to this parameter P.
- Show the Results
- Confirm the hypothesis.
- If there are results that are not according to the hypothesis, you have to be able to explain why this is the case!!!

9.3 Threats to Validity

Here you discuss any element of your experiments (usually depending on the setup), e.g., data sets used, parameters assessed, that might have implications for the validity of your results. For example, if you have not considered a specific range of parameter values, it is unclear how the algorithm will behave for these settings. If you have excluded certain data, this might have its implications for the generalizability of your results. In a way, this is the discussion section of your thesis.

10 Discussion

In this section, you discuss your general approach. What design decisions did you make, and, what are the implications of this? Maybe your approach computes an exact intermediary result, which is costly in terms of time. Does a suitable approximation approach exist? What are the implications for that approach. You can also dive deeper in very related work and explain why it does not work in certain cases.

11 Conclusion

Your conclusion actually follows the same structure as your abstract. However, you present what the reader has seen, again, using the same concepts as you have already used in your abstract.

Future Work. Furthermore, you present interesting directions for future work!

A Some additional stuff

Try to avoid appendices. If you can't, this is a place to show all kinds of results that are supportive, yet, not critical for your thesis.

Bibliography

- [1] B. R. Abarca Zúñiga, A. Yeshchenko, and H. van der Aa. “Version Clustering: A Top-Down Approach for Process Concept Drift Detection.” In: *Advanced Information Systems Engineering - 38th International Conference, CAiSE 2026, Verona, Italy, June 8-12, 2026, Proceedings, Part I*. Ed. by L. Fuentes, P. Plebani, C. Combi, and H. Reijers. Lecture Notes in Computer Science. Springer, 2026, pp. 61–77. DOI: 10.1007/978-3-032-28110-4_4.
- [2] J. N. Adams, C. Pitsch, T. Brockhoff, and W. M. P. van der Aalst. “An Experimental Evaluation of Process Concept Drift Detection.” In: *Proceedings of the VLDB Endowment* 16.8 (2023), pp. 1856–1869. DOI: 10.14778/3594512.3594517.
- [3] J. N. Adams, S. J. van Zelst, T. Rose, and W. M. P. van der Aalst. “Explainable Concept Drift in Process Mining.” In: *Information Systems* 114 (2023), p. 102177. DOI: 10.1016/j.is.2023.102177.
- [4] S. Barbon Junior, G. M. Tavares, V. G. Turrise da Costa, P. Ceravolo, and E. Damiani. “A Framework for Human-in-the-Loop Monitoring of Concept-Drift Detection in Event Log Stream.” In: *Companion of the The Web Conference 2018 on The Web Conference 2018, WWW 2018, Lyon, France, April 23-27, 2018*. ACM, 2018, pp. 319–326. DOI: 10.1145/3184558.3186343.
- [5] A. Berti, D. Kretschmann, and W. M. P. van der Aalst. *Identifying Execution States and Boundary Conditions in Object-Centric Event Logs: A Framework and Case Study*. TechRxiv preprint. 2025. DOI: 10.36227/techrxiv.176522567.74931712/v1.
- [6] R. P. J. C. Bose, W. M. P. van der Aalst, I. Žliobaitė, and M. Pechenizkiy. “Handling Concept Drift in Process Mining.” In: *Advanced Information Systems Engineering - 23rd International Conference, CAiSE 2011, London, UK, June 20-24, 2011. Proceedings*. Ed. by H. Mouratidis and C. Rolland. Vol. 6741. Lecture Notes in Computer Science. Springer, 2011, pp. 391–405. DOI: 10.1007/978-3-642-21640-4_30.
- [7] T. Brockhoff, M. S. Uysal, and W. M. P. van der Aalst. “Time-Aware Concept Drift Detection Using the Earth Mover’s Distance.” In: *2nd International Conference on Process Mining, ICPM 2020, Padua, Italy, October 5-8, 2020*. IEEE, 2020, pp. 33–40. DOI: 10.1109/ICPM49681.2020.00016.

- [8] J. Grimm, A. Kraus, and H. van der Aa. “CDLG: A Tool for the Generation of Event Logs with Concept Drifts.” In: *Proceedings of the Best Dissertation Award, Doctoral Consortium, and Demonstration & Resources Track at BPM 2022, Münster, Germany, September 11-16, 2022*. Ed. by C. Janiesch, C. Di Francescomarino, T. Grisold, A. Kumar, J. Mendling, B. Pentland, H. Reijers, M. Weske, and R. Winter. Vol. 3216. CEUR Workshop Proceedings. CEUR-WS.org, 2022, pp. 92–96. URL: https://ceur-ws.org/Vol-3216/paper_241.pdf.
- [9] S. Khraiweh, D. Dagtekin, and L. Pufahl. “CV-Driven Window Selection for Multi-Perspective Concept Drift Detection in Event Logs.” In: *Business Process Management Forum - BPM 2026 Forum, Toronto, ON, Canada, September 27 - October 2, 2026, Proceedings*. Ed. by F. M. Maggi, H. Leopold, F. Santoro, and A. Gal. Vol. 599. Lecture Notes in Business Information Processing. Springer, 2027, pp. 185–202. DOI: 10.1007/978-3-032-37877-4_11.
- [10] E. L. Klijn, F. Mannhardt, and D. Fahland. “Multi-Perspective Concept Drift Detection: Including the Actor Perspective.” In: *Advanced Information Systems Engineering - 36th International Conference, CAiSE 2024, Limassol, Cyprus, June 3-7, 2024, Proceedings*. Ed. by G. Guizzardi, F. Santoro, H. Mouratidis, and P. Soffer. Vol. 14663. Lecture Notes in Computer Science. Springer, 2024, pp. 141–157. DOI: 10.1007/978-3-031-61057-8_9.
- [11] A. Kraus, K. Amiri Elyasi, A. Rebmann, S. Hadian, and H. van der Aa. “A Framework for Steady-State Detection in Process Mining.” In: *Information Systems* 140 (2026), p. 102740. DOI: 10.1016/j.is.2026.102740.
- [12] A. Kraus and H. van der Aa. “Machine Learning-Based Detection of Concept Drift in Business Processes.” In: *Process Science* 2.1 (2025), p. 5. DOI: 10.1007/s44311-025-00012-w.
- [13] A. Kraus and H. van der Aa. “Comprehensive Characterization of Concept Drifts in Process Mining.” In: *Information Systems* 135 (2026), p. 102584. DOI: 10.1016/j.is.2025.102584.
- [14] A. Maaradji, M. Dumas, M. La Rosa, and A. Ostovar. “Detecting Sudden and Gradual Drifts in Business Processes from Execution Traces.” In: *IEEE Transactions on Knowledge and Data Engineering* 29.10 (2017), pp. 2140–2154. DOI: 10.1109/TKDE.2017.2720601.
- [15] A. Padella and F. Vinci. *Everything Real but the Drift: A Concept-Drift Benchmark Event Logs for Process Mining*. Zenodo dataset, version 1. 2026. DOI: 10.5281/zenodo.21005862.
- [16] D. M. V. Sato, S. C. De Freitas, J. P. Barddal, and E. E. Scalabrin. “A Survey on Concept Drift in Process Mining.” In: *ACM Computing Surveys* 54.9 (2022), 189:1–189:38. DOI: 10.1145/3472752.

- [17] A. Senderovich, C. Di Francescomarino, C. Ghidini, K. Jorbina, and F. M. Maggi. “Intra and Inter-Case Features in Predictive Process Monitoring: A Tale of Two Dimensions.” In: *Business Process Management - 15th International Conference, BPM 2017, Barcelona, Spain, September 10-15, 2017, Proceedings*. Ed. by J. Carmona, G. Engels, and A. Kumar. Vol. 10445. Lecture Notes in Computer Science. Springer, 2017, pp. 306–323. DOI: 10.1007/978-3-319-65000-5_18.
- [18] W. M. P. van der Aalst. *Process Mining: Data Science in Action*. 2nd ed. Springer, 2016. DOI: 10.1007/978-3-662-49851-4.
- [19] W. M. P. van der Aalst. *How to Write Beautiful Process-and-Data-Science Papers?* 2024. arXiv: 2203.09286 [cs.DL]. URL: <https://arxiv.org/abs/2203.09286>.
- [20] C. Zheng, L. Wen, and J. Wang. “Detecting Process Concept Drifts from Event Logs.” In: *On the Move to Meaningful Internet Systems. OTM 2017 Conferences - Confederated International Conferences: CoopIS, C&TC, and ODBASE 2017, Rhodes, Greece, October 23-27, 2017, Proceedings, Part I*. Ed. by H. Panetto, C. Debruyne, W. Gaaloul, M. Papazoglou, A. Paschke, C. A. Ardagna, and R. Meersman. Vol. 10573. Lecture Notes in Computer Science. Springer, 2017, pp. 524–542. DOI: 10.1007/978-3-319-69462-7_33.